

1. Good Evening
2. We begin at 9:07 pm
3. Topic - Transactions

## Agenda

1. Introduction
2. ACID
3. Start, Commit, Rollback, Demo
4. Isolation Levels - 4 levels
5. Demo
6. Deadlocks


Task [Money Transfer  $A \rightarrow B$ ]  
 Case1 [Crashes & Failures]  
 transferMoney (A, B, d)

1.	$x = \text{Read}(A)$	✓
2.	$x = x - d$	✓
3.	Write (A, x)	✓
4.	$x = \text{Read}(B)$	→ <u>crash/failure</u>
5.	$x = x + d$	
6.	Write (B, x)	

A (<sup>900</sup>~~1000~~)  $\xrightarrow{100}$  B (2000)  
 $x = 1000$

$x = x - 100$   
 Write(A, x)

## Case 2 [Concurrent Tasks]



- |   |  |
|---|--|
| <ol style="list-style-type: none"> <li>1. <math>x = \text{Read}(A)</math> ✓ 1000</li> <li>2. <math>x = x - 100</math> ✓ 900</li> <li>3. <math>\text{Write}(A, x)</math> ✓</li> <li>4. <math>x = \text{Read}(B)</math> ✓ 2000</li> <li>5. <math>x = x + 100</math> ✓ 2100</li> <li>6. <math>\text{Write}(B, x)</math> ✓</li> </ol> | <ol style="list-style-type: none"> <li>1. <math>x = \text{Read}(C)</math> ✓ 3000</li> <li>2. <math>x = x - 300</math> ✓ 2700</li> <li>3. <math>\text{Write}(C, x)</math> ✓</li> <li>4. <math>x = \text{Read}(B)</math> ✓ 2000</li> <li>5. <math>x = x + 300</math> ✓ 2300</li> <li>6. <math>\text{Write}(B, x)</math> ✓</li> </ol> |
|---|--|

Case 1 → Failures / Crashes can lead to data inconsistency.

Case 2 → Concurrent tasks can lead to data inconsistency.

## Transactions

↳ Set of operations to be run together.

ACID

A - Atomicity

.. ..  
C → Consistency  
I → Isolation  
D → Durability

A (Atomicity) → All or nothing  
↳ A transaction should be indivisible

C (Consistency) → State of data before & after the transaction should be consistent.

I (Isolation) → One transaction should not be impacted by other transactions in the system.

D (Durability) → Once the transaction finishes successfully the changes should not be lost

↙  
DBMS should make sure that changes got written to hard-disk at the end of transaction

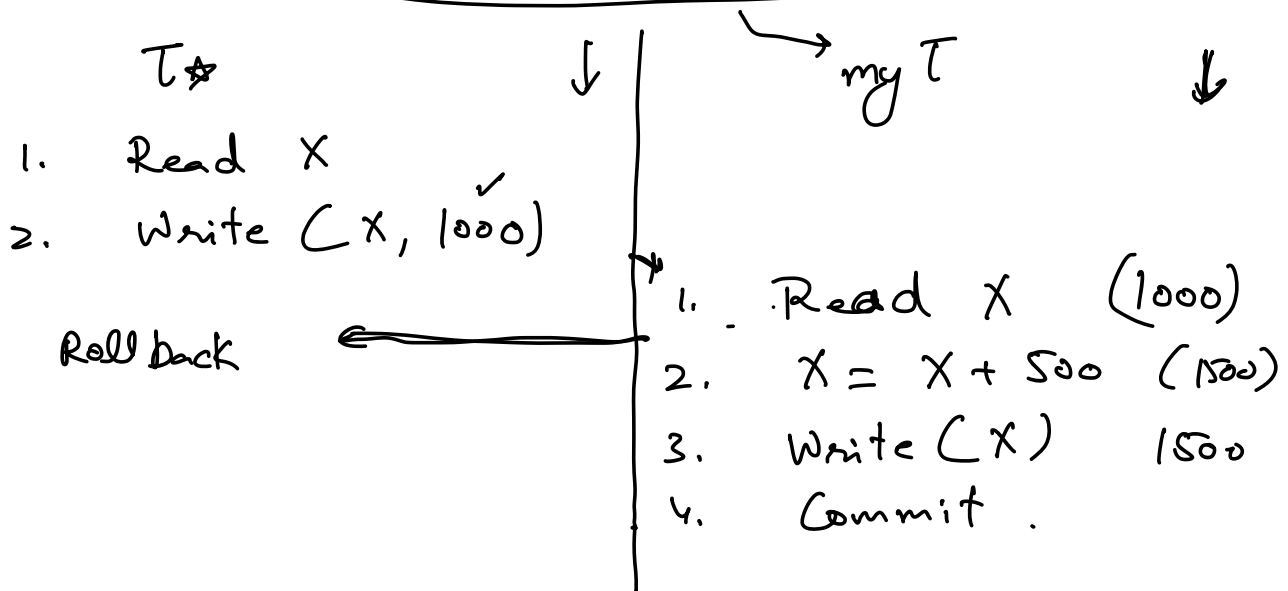
keywords  $\rightarrow$  Start, Commit, Rollback [Failure]

## $\rightarrow$ ISOLATION Levels & Concurrency

Concurrency vs Consistency

1. Read Uncommitted
2. Read Committed
3. Repeatable Read  $\rightarrow$  Default
4. Serializable

Read Uncommitted  $X = \overset{1500 \checkmark}{\cancel{1000}}$



Should we ever use Read Uncommitted?

$\hookrightarrow$  Financial Application X

Booking.com X  
 Social Media → Maybe if data is not critical

★ Pros of using read uncommitted  
 ↳ High Performance

Read Committed

→ Solves dirty read

$$X = \frac{2500}{2000} \checkmark$$

rt (X=1000)

my T (X+=500)

1. Read X 2000
2. Write (X, 1000) 1000 ✓
3. Rollback

1. Read X ✓ 2000
2. X = X + 500 ✓ 2500
3. Write (X) 2500
4. Commit ✓

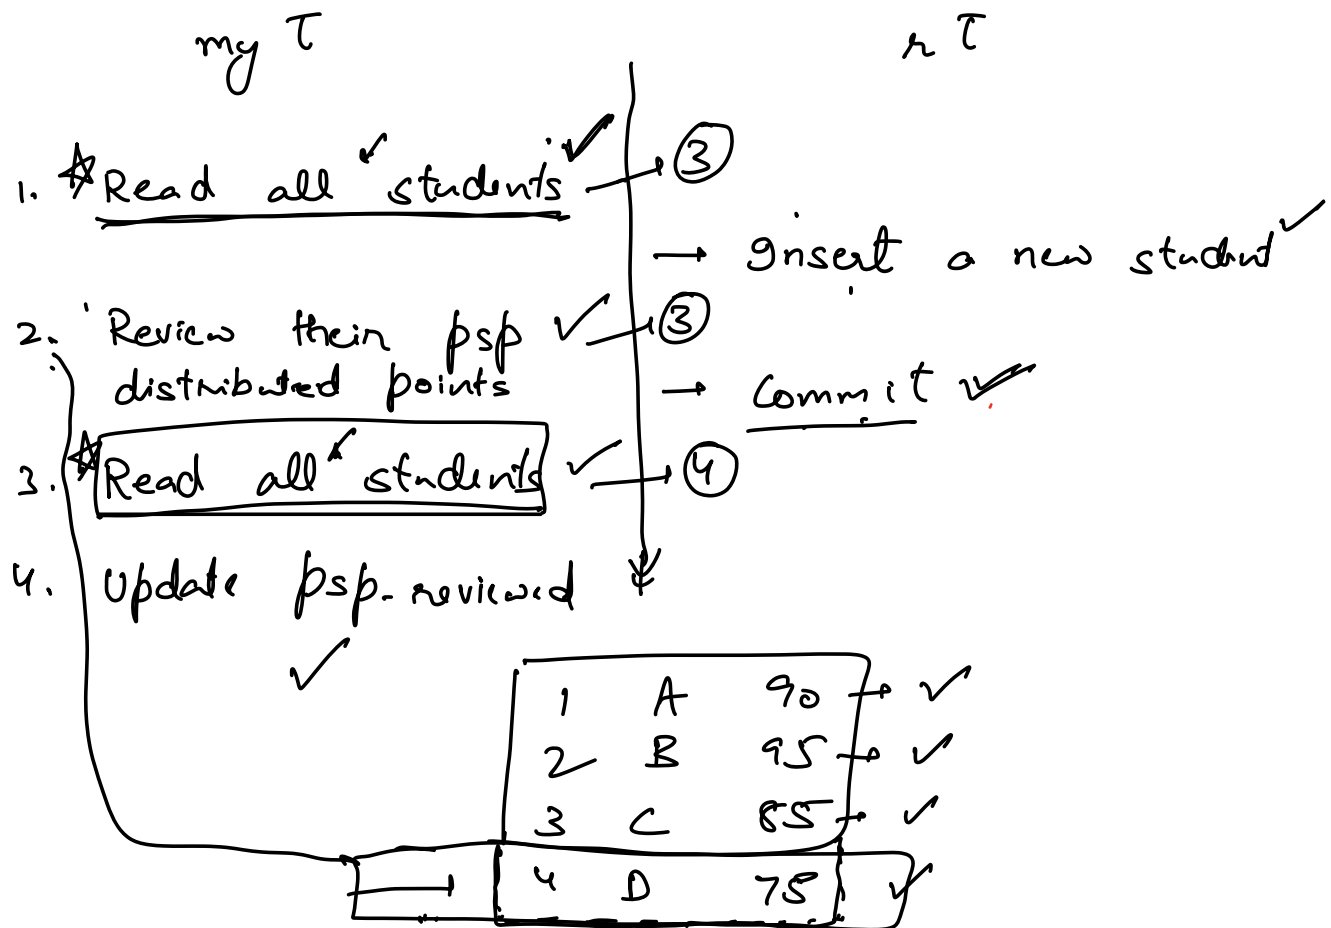
→ Problem with read committed [Non-repeatable  
 Read]

myT → Read all students

Increase the points with  $\geq 90$  psp

Update all students, that psp is reviewed

nt → Insert a new student



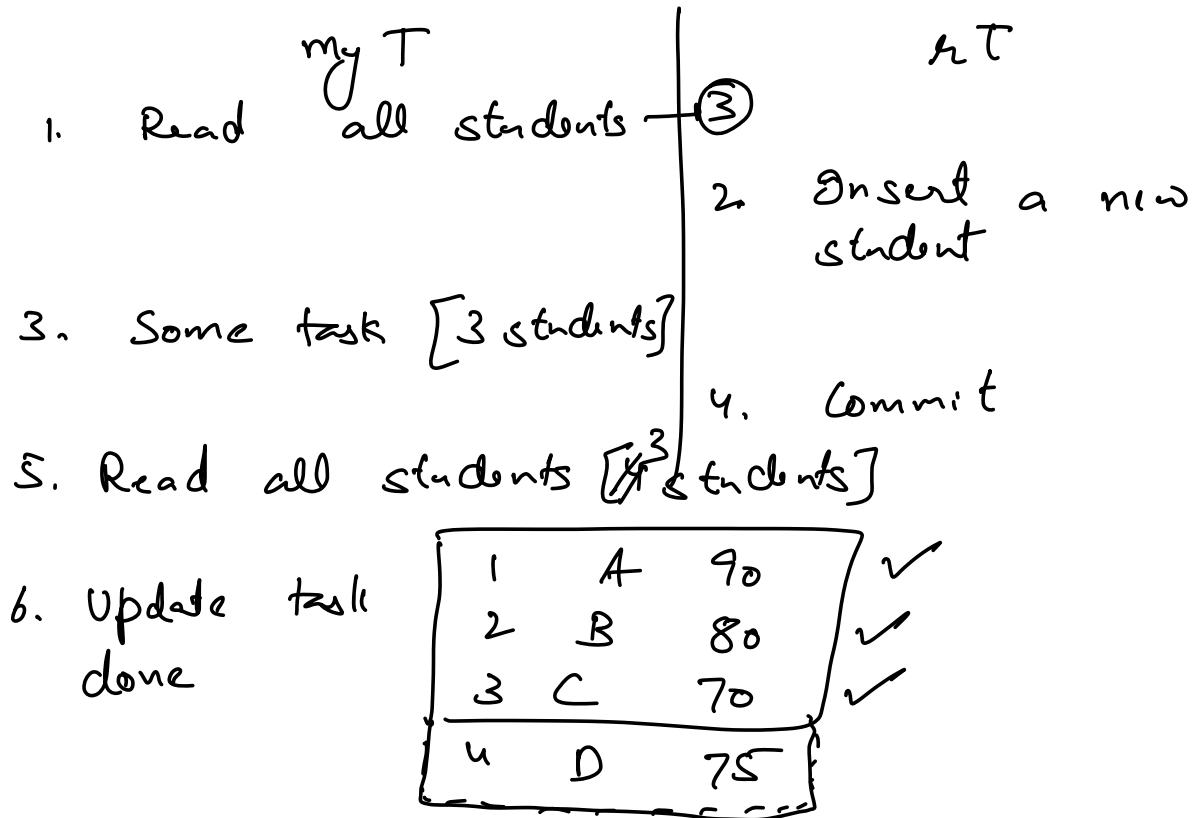
Read Uncommitted → dirty reads

Read Committed → ✓

Non-repeatable read

# REPEATABLE READ

→ Solve non-repeatable problem



1. Read uncommitted → Dirty Read ✓
2. Read Committed → Non-repeatable read ✓
3. Repeatable Read [Default]

(10:48 - 10:58)

Demo → RUC, RC, RR ✓  
 → Serializable  
 → Demo  
 → Deadlocks  
 → Demo

[Problem which even repeatable read  
 can't solve

READ SERIALIZE

A  $\xrightarrow{100}$  B  $\xleftarrow{200}$  C

my T (A → B)  $\begin{array}{r} 1000 \\ +200 \\ \hline 1100 \end{array}$

rT (C → B) ✓

✓ 1. x = read(B) 1000

1. x = read(B) 1000

2. x = x + 200 1200

3. Write (B, x)

4. Commit

2. x = x + 100 1100

3. Write (B, x)

4. Commit

A  $\xrightarrow{100}$  B  $\xleftarrow{200}$  C  $\begin{array}{r} 1000 \\ +200 \\ \hline 1200 \\ +100 \\ \hline 1300 \end{array}$

Read, Serializable

1. x = Read with lock (B)  
 Select Then update 1000



↓

1.  $x = \text{Read}(B)$  1200
2.  $x = x + 100$  (300)
3. Write
4. Commit

2.  $x = x + 200$  1200
3. Write(B, x)
4. Commit

- Dead locks
- 3 doubts → coalesce, set session, Doubt  
vs  
set trans
- Assignments → None for today